TABLE S6d

MODEL IDENTITIES AND INFERENCE SETTINGS

Model	Requested model ID	Transport	Reasoning	Decoding fields sent
GPT-6 Astra	gpt-6-astra; openai/gpt-6-astra	Batch	high	Combined output ceiling 32,768; seed 42; temperature/top-p omitted
Claude Opus 5.5	anthropic/claude-opus-5.5	Batch	high	Combined output ceiling 32,768; seed omitted; temperature/top-p omitted
Gemini 4 Argon	google/gemini-4-argon	Standard streaming	high	Combined output ceiling 943,718; seed 42; temperature/top-p omitted
DeepSeek V4.1 Flash 1017	deepseek/deepseek-v4.1-flash-1017	Standard streaming	high	Combined output ceiling 943,718; seed omitted; temperature/top-p omitted

Actual accepted settings are preserved per request with served model/provider, transport/recovery provenance and literal-body hashes. Exact accepted prior answers are reused once; interrupted or rejected attempts are separate physical records. One selected-means request and the separate eight-item Astra extension remain pending.